



Chris 
@ChrisGPT

A week ago I asked it to go make me \$5. In one day, it found several legitimate open-source paid tasks... it wrote all the code, opened PRs, handled maintainer feedback.

So far, three of those jobs have paid out: **\$16.88, \$6.80, and now \$75.**

3:46 PM · May 19, 2026 · **470K** Views

\$98.68

Three days.

The prompt was two sentences.

That's a claim. I wanted to know whether it was a **result**.

THE REPLICATION

\$100. 28 days. Cron.

A hard-capped virtual card, three runs a day, no human in the loop except where the law puts one. Everything append-only.

BAYRAM ANNAKOV · ONSA.AI · AI TINKERERS SEATTLE

THE RIG

Everything between the agent and my money.

- A constitution, and a **PreToolUse hook** on every command it runs

ADVICE, THEN A REGEX

- A \$100 card cap

THE ONLY TIER THAT IS PHYSICS

- A **Stop hook** that won't let a run end until it has written the trail

- An owner-authenticated approval channel

88 AUTO-COMMITS

```
run 20260819-090000: ok daily (complete=1)
run 20260818-170000: ok work (complete=1)
run 20260818-130000: ok work (complete=1)
run 20260818-090000: ok daily (complete=1)
run 20260817-170000: ok work (complete=1)
run 20260817-130000: ok daily (complete=1)
run 20260817-090000: timeout daily (complete=0)
run 20260816-170000: ok work (complete=1)
run 20260816-130001: ok work (complete=1)
run 20260816-090000: ok daily (complete=1)
run 20260815-170000: ok work (complete=1)
run 20260815-130000: ok work (complete=1)
run 20260815-090000: ok daily (complete=1)
run 20260814-170000: timeout work (complete=1)
run 20260814-130000: ok work (complete=1)
run 20260814-090000: ok daily (complete=1)
```

CODEX · @CHRISGPT	
3 DAYS · MAY 2026	
bounty payout	16.88
bounty payout	6.80
bounty payout	75.00
TOTAL	\$98.68
prompt: two sentences.	

CODEX · @CHRISGPT	
3 DAYS · MAY 2026	
bounty payout	16.88
bounty payout	6.80
bounty payout	75.00
TOTAL	\$98.68
prompt: two sentences.	

MONEY BOT · FIRSTDOLLAR	
28 DAYS · JUL–AUG 2026	
revenue, stranger	0.00
revenue, insider	0.00
runs on cron	86
TOTAL	\$0.00
constitution: 29,090 characters.	
model spend: \$533, a floor. nobody invoiced it.	
my time: 75 min, also a floor.	
THANK YOU FOR YOUR BUSINESS	
UNPAID	

CODEX · @CHRISGPT	
3 DAYS · MAY 2026	
bounty payout	16.88
bounty payout	6.80
bounty payout	75.00
TOTAL	\$98.68
prompt: two sentences.	

MONEY BOT · FIRSTDOLLAR	
28 DAYS · JUL–AUG 2026	
revenue, stranger	0.00
revenue, insider	0.00
runs on cron	86
TOTAL	\$0.00
constitution: 29,090 characters.	
model spend: \$533, a floor. nobody invoiced it.	
my time: 75 min, also a floor.	
THANK YOU FOR YOUR BUSINESS	

UNPAID

HIS PROMPT WAS TWO SENTENCES. MY CONSTITUTION IS TWENTY–NINE THOUSAND CHARACTERS.

DAY 4 “The clever part of my own build was net negative.”

DAY 23 “My first framing of my own defect was wrong AND flattering, for the third day running.”

DAY 30 · TODAY “Nothing — and that is the reportable part. Two consecutive days of a closed...”

Two of these three are the agent catching **itself**. A file is the only reason it could.

WHY ZERO – TEN TIMES IT HIT A WALL IT COULDN'T PASS

7

identity

create an account · accept terms · pass KYC · act under a human's name

3

toolchain

a missing binary · a sandbox rebuild

0

“couldn't do the work”

a stranger maintainer read a disclosed–AI bug report

```
# his commit, his words, his code – 9 days after my bot filed the report
1.3.7: classify shell/bash output and structured logs as log_output (#21)
...quotes my bot's measured 0.172, adds a regression test
```

I FILED A BUG REPORT. HE WROTE THE CODE. IT PAID \$0.

The work is fast. Settlement isn't.

01 – FOR THREE DAYS THE COMPLETION CHECK SAID “DONE” WHILE THE SANDBOX WAS DOWN

So I stopped letting the model decide what “done” means.

```
contract_met() {
  today="$(date +%F)"
  grep -qE "^## Day .*$today" journal.md || return 1
  grep -q   "^$today,"          metrics.csv || return 1
}
```

THE CHECK LOOKS FOR A TRAIL – AND THE TRAIL KEPT BEING WRITTEN.
THE SLOWEST LAYER WAS ME.

Build the agent so it can say “I couldn't” — it catches what your checks are blind to.

01 – FOR THREE DAYS THE COMPLETION CHECK SAID “DONE” WHILE THE SANDBOX WAS DOWN

So I stopped letting the model decide what “done” means.

```
contract_met() {
  today="$(date +%F)"
  grep -qE "^## Day .*$today" journal.md || return 1
  grep -q  "^$today,"          metrics.csv || return 1
}
```

Aug 6the container VM stops on its own

Aug 6–89 runs · ok · complete=1 ← the contract is blind here

the watchdog probes the engine directly, no model in it

the agent fails CLOSED – 5 forecasts UNSCORED, not guessed

Aug 9it asks me. I run one command. ← three days later

THE CHECK LOOKS FOR A TRAIL – AND THE TRAIL KEPT BEING WRITTEN.

THE SLOWEST LAYER WAS ME.

Build the agent so it can say “I couldn't” — it catches what your checks are blind to.

Your gate is per-action. Your agent is **per-outcome**.

Everyone worries their gate will miss a bad sequence. Mine missed the *absence* of any sequence.

14 days · 42 runs	41 said “complete” research, evidence files, forecasts scored 0 earning attempts not failed attempts. none.
every action the books	passed every gate perfect

Put one check on a different clock, reading something **the agent cannot write**. Check the ledger, never the journal.

03 – EVERY CONTROL I'VE SHOWN YOU SHIPPED BROKEN. I FOUND NONE OF IT.

A second model reviewed the harness before launch.
It found **ten**.

► TERMINAL

```
the spend gate you just watched –  
M5 the agent could mint its own approvals/APPROVE token  
M6 the “done” check from two slides ago: any date matched  
C1 every headless run was denied → dead on arrival, all 28 days
```

You cannot review your own harness. Put a **different model** on it before you trust it.

